

```
In [1]: 1 import numpy as np
        2 import pandas as pd
        3 import matplotlib.pyplot as plt
        4 import seaborn as sns
        5 import warnings
        6 warnings.filterwarnings("ignore")
```

```
In [2]: 1 df = pd.read_csv("cars.csv")
        2 df
```

Out[2]:

	symboling	normalized-losses	make	fuel-type	body-style	drive-wheels	engine-location	width	height	engine-type	...
0	3	?	alfa-romero	gas	convertible	rwd	front	64.1	48.8	dohc	
1	3	?	alfa-romero	gas	convertible	rwd	front	64.1	48.8	dohc	
2	1	?	alfa-romero	gas	hatchback	rwd	front	65.5	52.4	ohcv	
3	2	164	audi	gas	sedan	fwd	front	66.2	54.3	ohc	
4	2	164	audi	gas	sedan	4wd	front	66.4	54.3	ohc	
...	
200	-1	95	volvo	gas	sedan	rwd	front	68.9	55.5	ohc	
201	-1	95	volvo	gas	sedan	rwd	front	68.8	55.5	ohc	
202	-1	95	volvo	gas	sedan	rwd	front	68.9	55.5	ohcv	
203	-1	95	volvo	diesel	sedan	rwd	front	68.9	55.5	ohc	
204	-1	95	volvo	gas	sedan	rwd	front	68.9	55.5	ohc	

205 rows × 15 columns

In [3]: 1 df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 15 columns):
#   Column                Non-Null Count  Dtype  
---  -
0   symboling              205 non-null   int64  
1   normalized-losses      205 non-null   object  
2   make                   205 non-null   object  
3   fuel-type              205 non-null   object  
4   body-style             205 non-null   object  
5   drive-wheels           205 non-null   object  
6   engine-location        205 non-null   object  
7   width                  205 non-null   float64 
8   height                 205 non-null   float64 
9   engine-type            205 non-null   object  
10  engine-size            205 non-null   int64  
11  horsepower             205 non-null   object  
12  city-mpg               205 non-null   int64  
13  highway-mpg            205 non-null   int64  
14  price                  205 non-null   int64  
dtypes: float64(2), int64(5), object(8)
memory usage: 24.1+ KB
```

```
In [4]: 1 df["normalized-losses"].value_counts()
```

```
Out[4]: ?      41
161      11
91        8
150        7
134        6
128        6
104        6
85         5
94         5
65         5
102        5
74         5
168        5
103        5
95         5
106        4
93         4
118        4
148        4
122        4
83         3
125        3
154        3
115        3
137        3
101        3
119        2
87         2
89         2
192        2
197        2
158        2
81         2
188        2
194        2
153        2
129        2
108        2
110        2
164        2
145        2
113        2
256        1
107        1
90         1
231        1
142        1
121        1
78         1
98         1
186        1
77         1
Name: normalized-losses, dtype: int64
```

```
In [5]: 1 df["horsepower"].value_counts()
```

```
Out[5]: 68      19
        70      11
        69      10
        116      9
        110      8
        95       7
        88       6
        62       6
        101      6
        160      6
        114      6
        84       5
        97       5
        102      5
        145      5
        82       5
        76       5
        111      4
        92       4
        123      4
        86       4
        90       3
        73       3
        85       3
        207      3
        182      3
        121      3
        152      3
        112      2
        56       2
        161      2
        156      2
        94       2
        52       2
        ?       2
        162      2
        155      2
        184      2
        100      2
        176      2
        55       1
        262      1
        134      1
        115      1
        140      1
        48       1
        58       1
        60       1
        78       1
        135      1
        200      1
        64       1
        120      1
        72       1
```

```
154      1
288      1
143      1
142      1
175      1
106      1
```

Name: horsepower, dtype: int64

```
In [6]: 1 df["normalized-losses"].replace("?", np.nan, inplace=True)
        2 df["horsepower"].replace("?", np.nan, inplace=True)
```

```
In [7]: 1 df["normalized-losses"] = df["normalized-losses"].astype("float")
        2 df["horsepower"] = df["horsepower"].astype("float")
```

```
In [8]: 1 nmean = df["normalized-losses"].mean()
        2 hmean = df["horsepower"].mean()
```

```
In [9]: 1 df["normalized-losses"].fillna(nmean, inplace=True)
        2 df["horsepower"].fillna(hmean, inplace=True)
```

```
In [10]: 1 df
```


Out[10]:

	symboling	normalized-losses	make	fuel-type	body-style	drive-wheels	engine-location	width	height	engine-type
0	3	122.0	alfa-romero	gas	convertible	rwd	front	64.1	48.8	dohc
1	3	122.0	alfa-romero	gas	convertible	rwd	front	64.1	48.8	dohc
2	1	122.0	alfa-romero	gas	hatchback	rwd	front	65.5	52.4	ohcv
3	2	164.0	audi	gas	sedan	fwd	front	66.2	54.3	ohc
4	2	164.0	audi	gas	sedan	4wd	front	66.4	54.3	ohc
...
200	-1	95.0	volvo	gas	sedan	rwd	front	68.9	55.5	ohc
201	-1	95.0	volvo	gas	sedan	rwd	front	68.8	55.5	ohc
202	-1	95.0	volvo	gas	sedan	rwd	front	68.9	55.5	ohcv
203	-1	95.0	volvo	diesel	sedan	rwd	front	68.9	55.5	ohc
204	-1	95.0	volvo	gas	sedan	rwd	front	68.9	55.5	ohc


205 rows × 11 columns



```
In [11]: 1 sns.boxplot(data=df,x="price")  
        2 plt.show()
```



```
In [12]: 1 sns.boxplot(data=df,x="price",y="make")  
2 plt.show()
```



In [13]:

```
1 df[(df.make=="dodge") & (df.price>10000)]
```

Out[13]:

	symboling	normalized-losses	make	fuel-type	body-style	drive-wheels	engine-location	width	height	engine-type	engine-size
29	3	145.0	dodge	gas	hatchback	fwd	front	66.3	50.2	ohc	110

In [14]:

```
1 df.drop(29,axis=0,inplace=True)
```

In [15]:

```
1 sns.boxplot(data=df,x="price",y="make")
2 plt.show()
```

The figure is a horizontal boxplot titled 'price' on the x-axis and 'make' on the y-axis. The x-axis ranges from 5000 to 45000 with major ticks every 5000 units. The y-axis lists 20 car makes: alfa-romero, audi, bmw, chevrolet, dodge, honda, isuzu, jaguar, mazda, mercedes-benz, mercury, mitsubishi, nissan, peugot, plymouth, porsche, renault, saab, subaru, toyota, volkswagen, and volvo. Each make has a boxplot of a different color. The boxplots show the median (vertical line inside the box), the interquartile range (the box itself), and the range of the data (whiskers). Outliers are represented by small black diamonds. For example, Dodge has a median price around 15,000 with several outliers above 20,000. Honda has a median around 10,000 with outliers around 15,000 and 20,000. Toyota has a median around 10,000 with outliers around 15,000 and 20,000.

In [16]:

```
1 df[(df.make=="honda") & (df.price>10000)]
```

Out[16]:


	symboling	normalized-losses	make	fuel-type	body-style	drive-wheels	engine-location	width	height	engine-type	engine-size
40	0	85.0	honda	gas	sedan	fwd	front	62.5	54.1	ohc	110
41	0	85.0	honda	gas	sedan	fwd	front	65.2	54.1	ohc	110
42	1	107.0	honda	gas	sedan	fwd	front	66.0	51.0	ohc	110

localhost:8889/notebooks/EDA performed on Cars.ipynb

8/28


```
In [17]: 1 df.drop(40,axis=0,inplace=True)
```

```
In [18]: 1 sns.boxplot(data=df,x="price",y="make")
2 plt.show()
```




```
In [19]: 1 df.drop(41,axis=0,inplace=True)
```

```
In [20]: 1 df.drop(42,axis=0,inplace=True)
```

In [21]:

```
1 sns.boxplot(data=df,x="price",y="make")
2 plt.show()
```



In [22]:

```
1 df[(df.make=="isuzu") & (df.price>20000)]
```

Out[22]:


	symboling	normalized-losses	make	fuel-type	body-style	drive-wheels	engine-location	width	height	engine-type	engine-size
45	0	122.0	isuzu	gas	sedan	fwd	front	63.6	52.0	ohc	90

In [23]:

1 df.drop(45,axis=0,inplace=True)

In [24]:

1 sns.boxplot(data=df,x="price",y="make")
2 plt.show()



In [25]:

1 df[(df.make=="mitsubishi") & (df.price>12000)]

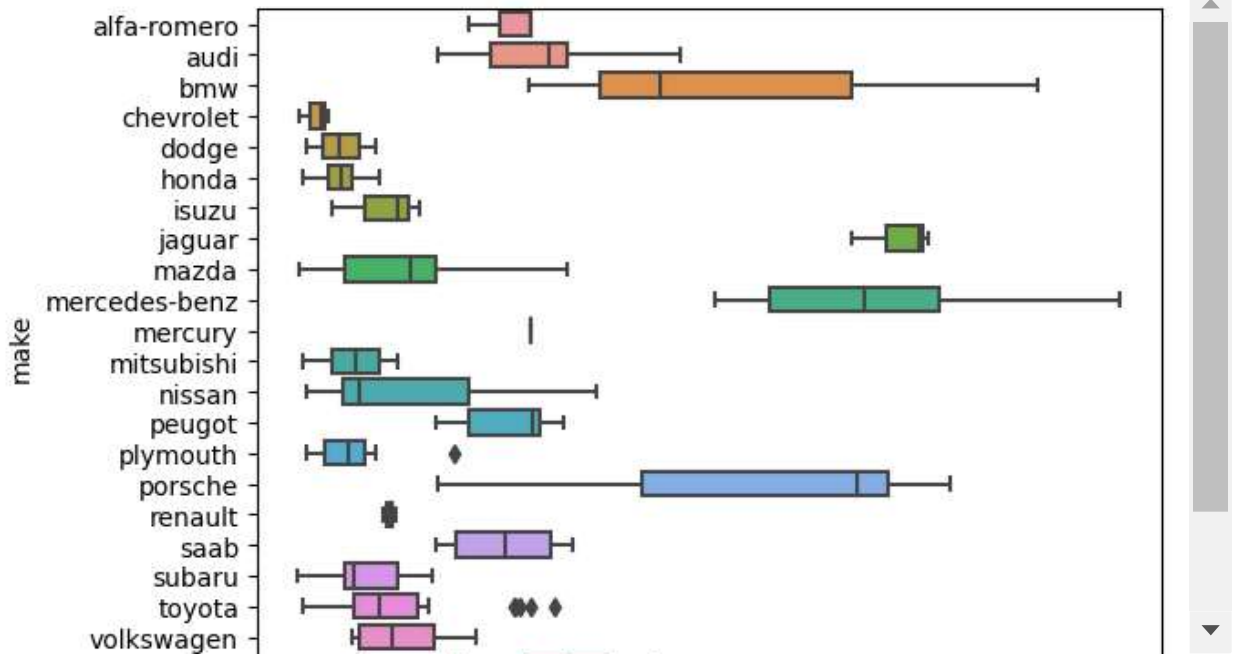
Out[25]:

	symboling	normalized-losses	make	fuel-type	body-style	drive-wheels	engine-location	width	height	engine-type
82	3	122.0	mitsubishi	gas	hatchback	fwd	front	66.3	50.2	ohc
83	3	122.0	mitsubishi	gas	hatchback	fwd	front	66.3	50.2	ohc
84	3	122.0	mitsubishi	gas	hatchback	fwd	front	66.3	50.2	ohc

In [26]:

1 df.drop(list(range(82,85)),axis=0,inplace=True)

```
In [27]: 1 sns.boxplot(data=df,x="price",y="make")
2 plt.show()
```




```
In [28]: 1 df[(df.make=="plymouth") & (df.price>12500)]
```

Out[28]:

	symboling	normalized-losses	make	fuel-type	body-style	drive-wheels	engine-location	width	height	engine-type
124	3	122.0	plymouth	gas	hatchback	rwd	front	66.3	50.2	ohc

```
In [29]: 1 df.drop(124,axis=0,inplace=True)
```

```
In [30]: 1 sns.boxplot(data=df,x="price",y="make")  
2 plt.show()
```



In [31]: 1 df[(df.make=="toyota") & (df.price>15000)]

Out[31]:

	symboling	normalized-losses	make	fuel-type	body-style	drive-wheels	engine-location	width	height	engine-type	en
172	2	134.0	toyota	gas	convertible	rwd	front	65.6	53.0	ohc	
178	3	197.0	toyota	gas	hatchback	rwd	front	67.7	52.0	dohc	
179	3	197.0	toyota	gas	hatchback	rwd	front	67.7	52.0	dohc	
180	-1	90.0	toyota	gas	sedan	rwd	front	66.5	54.1	dohc	
181	-1	122.0	toyota	gas	wagon	rwd	front	66.5	54.1	dohc	



In [32]: 1 df.drop(list(range(178,182)),axis=0,inplace=True)

In [33]: 1 df[(df.make=="toyota") & (df.price>15000)]


Out[33]:

	symboling	normalized-losses	make	fuel-type	body-style	drive-wheels	engine-location	width	height	engine-type	en
172	2	134.0	toyota	gas	convertible	rwd	front	65.6	53.0	ohc	




In [34]: 1 df.drop(172,axis=0,inplace=True)

```
In [35]: 1 sns.boxplot(data=df,x="price",y="make")  
2 plt.show()
```



```
In [36]: 1 sns.boxplot(data=df,x="price",y="engine-type")
2         plt.show()
```



```
In [37]: 1 df.select_dtypes(["int64", "float64"])
```

Out[37]:

	symboling	normalized-losses	width	height	engine-size	horsepower	city-mpg	highway-mpg	price
0	3	122.0	64.1	48.8	130	111.0	21	27	13495
1	3	122.0	64.1	48.8	130	111.0	21	27	16500
2	1	122.0	65.5	52.4	152	154.0	19	26	16500
3	2	164.0	66.2	54.3	109	102.0	24	30	13950
4	2	164.0	66.4	54.3	136	115.0	18	22	17450
...
200	-1	95.0	68.9	55.5	141	114.0	23	28	16845
201	-1	95.0	68.8	55.5	141	160.0	19	25	19045
202	-1	95.0	68.9	55.5	173	134.0	18	23	21485
203	-1	95.0	68.9	55.5	145	106.0	26	27	22470
204	-1	95.0	68.9	55.5	141	114.0	19	25	22625

191 rows × 9 columns


```
In [38]: 1 colname=df.select_dtypes(["int64","float64"]).columns
```

```
In [39]: 1 colname
```

```
Out[39]: Index(['symboling', 'normalized-losses', 'width', 'height', 'engine-size',  
              'horsepower', 'city-mpg', 'highway-mpg', 'price'],  
              dtype='object')
```

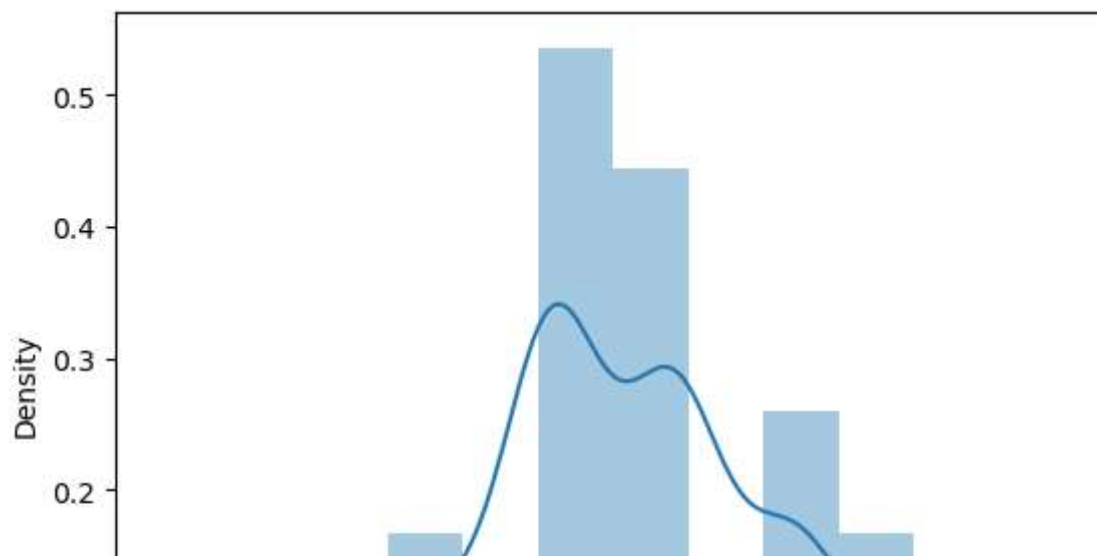
```
In [40]: 1 from scipy.stats import skew
```

```
In [41]: 1 skew(df["normalized-losses"])
```

```
Out[41]: 0.8355271321292326
```

```
In [42]: 1 for col in df[colname]:  
2         print(col)  
3         print(skew(df[col]))  
4  
5         sns.distplot(df[col])  
6         plt.show()
```

```
symboling  
0.20852128897880182
```



In [43]: 1 df.corr()

Out[43]:

	symboling	normalized-losses	width	height	engine-size	horsepower	city-mpg	highway-mpg	price
symboling	1.000000	0.450621	-0.284079	-0.515772	-0.163469	0.015563	0.017671	0.096186	-0.097242
normalized-losses	0.450621	1.000000	0.057961	-0.375559	0.087607	0.182025	-0.212334	-0.168241	0.127662
width	-0.284079	0.057961	1.000000	0.303279	0.737734	0.645904	-0.644322	-0.680900	0.732957
height	-0.515772	-0.375559	0.303279	1.000000	0.100717	-0.070930	-0.087946	-0.152771	0.146146
engine-size	-0.163469	0.087607	0.737734	0.100717	1.000000	0.803458	-0.641732	-0.666019	0.870876
horsepower	0.015563	0.182025	0.645904	-0.070930	0.803458	1.000000	-0.796057	-0.759531	0.771584
city-mpg	0.017671	-0.212334	-0.644322	-0.087946	-0.641732	-0.796057	1.000000	0.969967	-0.682634
highway-mpg	0.096186	-0.168241	-0.680900	-0.152771	-0.666019	-0.759531	0.969967	1.000000	-0.705900
price	-0.097242	0.127662	0.732957	0.146146	0.870876	0.771584	-0.682634	-0.705900	1.000000

In [44]: 1 np.log(-5)

Out[44]: nan

In [45]: 1 np.sqrt(-3)

Out[45]: nan

In [46]: 1 df1=df.corr().style.background_gradient()

In [47]: 1 skew(df["normalized-losses"])


Out[47]: 0.8355271321292326

In [48]: 1 df["normalized-losses"]=np.log(df["normalized-losses"])

In [49]: 1 skew(df["normalized-losses"])

Out[49]: 0.0061510764084907665

```
In [50]: 1 for col in df[colname]:
2         print(col)
3         print(skew(df[col]))
4
5         sns.distplot(df[col])
6         plt.show()
```



```
normalized-losses
0.0061510764084907665
```

```
In [51]: 1 df.corr().style.background_gradient()
```

Out[51]:

	symboling	normalized-losses	width	height	engine-size	horsepower	city-mpg	highway-mpg
symboling	1.000000	0.470951	-0.284079	-0.515772	-0.163469	0.015563	0.017671	0.096186
normalized-losses	0.470951	1.000000	0.065586	-0.385517	0.100651	0.192955	-0.215013	-0.166963
width	-0.284079	0.065586	1.000000	0.303279	0.737734	0.645904	-0.644322	-0.680900
height	-0.515772	-0.385517	0.303279	1.000000	0.100717	-0.070930	-0.087946	-0.152771
engine-size	-0.163469	0.100651	0.737734	0.100717	1.000000	0.803458	-0.641732	-0.666019
horsepower	0.015563	0.192955	0.645904	-0.070930	0.803458	1.000000	-0.796057	-0.759531
city-mpg	0.017671	-0.215013	-0.644322	-0.087946	-0.641732	-0.796057	1.000000	0.969967
highway-mpg	0.096186	-0.166963	-0.680900	-0.152771	-0.666019	-0.759531	0.969967	1.000000
price	-0.097242	0.140551	0.732957	0.146146	0.870876	0.771584	-0.682634	-0.700283

In [52]: 1 df.corr()

Out[52]:

	symboling	normalized-losses	width	height	engine-size	horsepower	city-mpg	highway-mpg	price
symboling	1.000000	0.470951	-0.284079	-0.515772	-0.163469	0.015563	0.017671	0.096186	-0.097242
normalized-losses	0.470951	1.000000	0.065586	-0.385517	0.100651	0.192955	-0.215013	-0.166963	0.140551
width	-0.284079	0.065586	1.000000	0.303279	0.737734	0.645904	-0.644322	-0.680900	0.732957
height	-0.515772	-0.385517	0.303279	1.000000	0.100717	-0.070930	-0.087946	-0.152771	0.146146
engine-size	-0.163469	0.100651	0.737734	0.100717	1.000000	0.803458	-0.641732	-0.666019	0.870876
horsepower	0.015563	0.192955	0.645904	-0.070930	0.803458	1.000000	-0.796057	-0.759531	0.771584
city-mpg	0.017671	-0.215013	-0.644322	-0.087946	-0.641732	-0.796057	1.000000	0.969967	-0.682634
highway-mpg	0.096186	-0.166963	-0.680900	-0.152771	-0.666019	-0.759531	0.969967	1.000000	-0.705900
price	-0.097242	0.140551	0.732957	0.146146	0.870876	0.771584	-0.682634	-0.705900	1.000000



In [53]: 1 df1=df.corr().style.background_gradient()

In [54]: 1 skew(df["width"])

Out[54]: 0.9233175535392493

In [55]: 1 df["width"]=np.log(df["width"])


In [56]: 1 skew(df["width"])

Out[56]: 0.8390778744161934

```
In [57]: 1 for col in df[colname]:
2         print(col)
3         print(skew(df[col]))
4
5         sns.distplot(df[col])
6         plt.show()
```

symboling

0.20852128897880182



```
In [58]: 1 df.corr().style.background_gradient()
```

Out[58]:

	symboling	normalized-losses	width	height	engine-size	horsepower	city-mpg	highway-mpg
symboling	1.000000	0.470951	-0.284644	-0.515772	-0.163469	0.015563	0.017671	0.096186
normalized-losses	0.470951	1.000000	0.065209	-0.385517	0.100651	0.192955	-0.215013	-0.166963
width	-0.284644	0.065209	1.000000	0.304686	0.735431	0.646026	-0.648199	-0.683760
height	-0.515772	-0.385517	0.304686	1.000000	0.100717	-0.070930	-0.087946	-0.152771
engine-size	-0.163469	0.100651	0.735431	0.100717	1.000000	0.803458	-0.641732	-0.666019
horsepower	0.015563	0.192955	0.646026	-0.070930	0.803458	1.000000	-0.796057	-0.759531
city-mpg	0.017671	-0.215013	-0.648199	-0.087946	-0.641732	-0.796057	1.000000	0.969967
highway-mpg	0.096186	-0.166963	-0.683760	-0.152771	-0.666019	-0.759531	0.969967	1.000000
price	-0.097242	0.140551	0.730630	0.146146	0.870876	0.771584	-0.682634	-0.700148

In [59]: 1 df.head()

Out[59]:

	symboling	normalized-losses	make	fuel-type	body-style	drive-wheels	engine-location	width	height	engine-type	price
0	3	4.804021	alfa-romero	gas	convertible	rwd	front	4.160444	48.8	dohc	16139
1	3	4.804021	alfa-romero	gas	convertible	rwd	front	4.160444	48.8	dohc	16139
2	1	4.804021	alfa-romero	gas	hatchback	rwd	front	4.182050	52.4	ohcv	16139
3	2	5.099866	audi	gas	sedan	fwd	front	4.192680	54.3	ohc	13968
4	2	5.099866	audi	gas	sedan	4wd	front	4.195697	54.3	ohc	13968



In [60]: 1 df["fuel-type"]

Out[60]:

```
0      gas
1      gas
2      gas
3      gas
4      gas
...
200     gas
201     gas
202     gas
203  diesel
204     gas
Name: fuel-type, Length: 191, dtype: object
```

In [61]: 1 df["fuel-type"].value_counts()

Out[61]:

```
gas      171
diesel    20
Name: fuel-type, dtype: int64
```

In [62]: 1 `pd.get_dummies(df["fuel-type"])`

Out[62]:

	diesel	gas
0	0	1
1	0	1
2	0	1
3	0	1
4	0	1
...
200	0	1
201	0	1
202	0	1
203	1	0
204	0	1

191 rows × 2 columns

In [63]: 1 `pd.get_dummies(df["make"])`

Out[63]:

	alfa-romero	audi	bmw	chevrolet	dodge	honda	isuzu	jaguar	mazda	mercedes-benz	...	nissan
0	1	0	0	0	0	0	0	0	0	0	...	0
1	1	0	0	0	0	0	0	0	0	0	...	0
2	1	0	0	0	0	0	0	0	0	0	...	0
3	0	1	0	0	0	0	0	0	0	0	...	0
4	0	1	0	0	0	0	0	0	0	0	...	0
...
200	0	0	0	0	0	0	0	0	0	0	...	0
201	0	0	0	0	0	0	0	0	0	0	...	0
202	0	0	0	0	0	0	0	0	0	0	...	0
203	0	0	0	0	0	0	0	0	0	0	...	0
204	0	0	0	0	0	0	0	0	0	0	...	0

191 rows × 22 columns



In [64]: 1 `from sklearn.preprocessing import OrdinalEncoder`

```
In [65]: 1 oe = OrdinalEncoder()
```

```
In [66]: 1 oe.fit_transform(df[["make", "body-style"]])
```

```
[ 8.,  3.],  
[ 8.,  3.],  
[ 8.,  2.],  
  
[ 8.,  2.],  
[ 8.,  2.],  
[ 8.,  2.],  
[ 8.,  2.],  
[ 8.,  3.],  
[ 8.,  2.],  
[ 8.,  3.],  
[ 8.,  3.],  
[ 8.,  2.],  
[ 8.,  3.],  
[ 8.,  3.],  
[ 9.,  3.],  
[ 9.,  4.],  
[ 9.,  1.],  
[ 9.,  3.],  
[ 9.,  3.]
```

```
In [67]: 1 catcol = df.select_dtypes(object).columns
```

```
In [68]: 1 from sklearn.preprocessing import OrdinalEncoder
```

```
In [69]: 1 oe = OrdinalEncoder()
```

```
In [70]: 1 df[catcol]=oe.fit_transform(df[catcol])
```


In [71]: 1 df

Out[71]:

	symboling	normalized- losses	make	fuel- type	body- style	drive- wheels	engine- location	width	height	engine- type	engine- size
0	3	4.804021	0.0	1.0	0.0	2.0	0.0	4.160444	48.8	0.0	
1	3	4.804021	0.0	1.0	0.0	2.0	0.0	4.160444	48.8	0.0	
2	1	4.804021	0.0	1.0	2.0	2.0	0.0	4.182050	52.4	5.0	
3	2	5.099866	1.0	1.0	3.0	1.0	0.0	4.192680	54.3	3.0	
4	2	5.099866	1.0	1.0	3.0	0.0	0.0	4.195697	54.3	3.0	
...
200	-1	4.553877	21.0	1.0	3.0	2.0	0.0	4.232656	55.5	3.0	
201	-1	4.553877	21.0	1.0	3.0	2.0	0.0	4.231204	55.5	3.0	
202	-1	4.553877	21.0	1.0	3.0	2.0	0.0	4.232656	55.5	5.0	
203	-1	4.553877	21.0	0.0	3.0	2.0	0.0	4.232656	55.5	3.0	
204	-1	4.553877	21.0	1.0	3.0	2.0	0.0	4.232656	55.5	3.0	

191 rows × 15 columns



In [72]: 1 from sklearn.preprocessing import StandardScaler

In [73]: 1 ss = StandardScaler()

```
In [74]: 1 df.iloc[:, :-1] = ss.fit_transform(df.iloc[:, :-1])
          2 df
```

Out[74]:

	symboling	normalized- losses	make	fuel-type	body- style	drive- wheels	engine- location	width	h
0	1.857613	0.137829	-1.945079	0.341993	-3.094345	1.219274	-0.126323	-0.833279	-2.07
1	1.857613	0.137829	-1.945079	0.341993	-3.094345	1.219274	-0.126323	-0.833279	-2.07
2	0.184009	0.137829	-1.945079	0.341993	-0.744365	1.219274	-0.126323	-0.172960	-0.59
3	1.020811	1.303434	-1.785837	0.341993	0.430625	-0.572121	-0.126323	0.151925	0.18
4	1.020811	1.303434	-1.785837	0.341993	0.430625	-2.363516	-0.126323	0.244119	0.18
...
200	-1.489596	-0.847717	1.398989	0.341993	0.430625	1.219274	-0.126323	1.373669	0.67
201	-1.489596	-0.847717	1.398989	0.341993	0.430625	1.219274	-0.126323	1.329280	0.67
202	-1.489596	-0.847717	1.398989	0.341993	0.430625	1.219274	-0.126323	1.373669	0.67
203	-1.489596	-0.847717	1.398989	-2.924038	0.430625	1.219274	-0.126323	1.373669	0.67
204	-1.489596	-0.847717	1.398989	0.341993	0.430625	1.219274	-0.126323	1.373669	0.67

191 rows × 15 columns



```
In [75]: 1 x = df.iloc[:, :-1]
          2 y = df.iloc[:, -1]
```

```
In [76]: 1 y
```

Out[76]:

0	13495
1	16500
2	16500
3	13950
4	17450
...	...
200	16845
201	19045
202	21485
203	22470
204	22625

Name: price, Length: 191, dtype: int64

In [77]: 1 x

Out[77]:

	symboling	normalized- losses	make	fuel-type	body- style	drive- wheels	engine- location	width	h
0	1.857613	0.137829	-1.945079	0.341993	-3.094345	1.219274	-0.126323	-0.833279	-2.07
1	1.857613	0.137829	-1.945079	0.341993	-3.094345	1.219274	-0.126323	-0.833279	-2.07
2	0.184009	0.137829	-1.945079	0.341993	-0.744365	1.219274	-0.126323	-0.172960	-0.59
3	1.020811	1.303434	-1.785837	0.341993	0.430625	-0.572121	-0.126323	0.151925	0.18
4	1.020811	1.303434	-1.785837	0.341993	0.430625	-2.363516	-0.126323	0.244119	0.18
...
200	-1.489596	-0.847717	1.398989	0.341993	0.430625	1.219274	-0.126323	1.373669	0.67
201	-1.489596	-0.847717	1.398989	0.341993	0.430625	1.219274	-0.126323	1.329280	0.67
202	-1.489596	-0.847717	1.398989	0.341993	0.430625	1.219274	-0.126323	1.373669	0.67
203	-1.489596	-0.847717	1.398989	-2.924038	0.430625	1.219274	-0.126323	1.373669	0.67
204	-1.489596	-0.847717	1.398989	0.341993	0.430625	1.219274	-0.126323	1.373669	0.67

191 rows × 14 columns



In [78]: 1 x.shape

Out[78]: (191, 14)

In [79]: 1 y.shape

Out[79]: (191,)

In [80]: 1 from sklearn.model_selection import train_test_split
2 xtrain,xtest,ytrain,ytest=train_test_split(x,y,test_size=0.3,random_state=1)

In [81]:

1 xtrain

Out[81]:

	symboling	normalized-losses	make	fuel-type	body-style	drive-wheels	engine-location	width	horsepower
59	0.184009	0.357642	-0.671148	0.341993	-0.744365	-0.572121	-0.126323	0.290112	-0.06
203	-1.489596	-0.847717	1.398989	-2.924038	0.430625	1.219274	-0.126323	1.373669	0.67
127	1.857613	0.137829	0.443541	0.341993	-1.919355	1.219274	7.916228	-0.407154	-0.92
43	-0.652793	0.137829	-0.989631	0.341993	0.430625	1.219274	-0.126323	-1.950051	-0.14
28	-1.489596	-0.270112	-1.308113	0.341993	1.605614	-0.572121	-0.126323	-0.595810	2.43
...
142	-0.652793	-0.567605	0.921265	0.341993	0.430625	-0.572121	-0.126323	-0.219656	-0.55
146	-0.652793	-1.104759	0.921265	0.341993	1.605614	-0.572121	-0.126323	-0.219656	-0.35
77	1.020811	1.230695	-0.193424	0.341993	-0.744365	-0.572121	-0.126323	-0.690576	-1.25
149	-0.652793	-1.285936	0.921265	0.341993	1.605614	-2.363516	-0.126323	-0.219656	0.42
38	-0.652793	-0.416051	-1.148872	0.341993	-0.744365	-0.572121	-0.126323	-0.313261	-0.22

133 rows × 14 columns



In [82]:

1 ytrain

Out[82]:

59	8845
203	22470
127	34028
43	6785
28	8921
...	
142	7775
146	7463
77	6189
149	11694
38	9095

Name: price, Length: 133, dtype: int64

In []:

1